

ASSIGNMENT 3

Group 2:-

- 1)Ritish Bansal 190101076
- 2)Suryansh Singh 190101089
- 3)Anant Shankhdhar 190101011
- 4)Mayank Chandak 190101052

Part A

Lazy Memory allocation

For this part of assignment mainly 3 files were modified:

1) sysproc.c - The **sbrk()** function was changed so that physical memory was not allocated before it was needed. When any process needs physical memory then only it is allocated. We commented out the function which allocates physical memory and also increased **myproc()->sz** by **n**.

```
39 int
40 sys_getpid(void)
41 {
42     return myproc()->pid;
43 }
44
45 int
46 sys_sbrk(void)
47 {
48     int addr;
49     int n;
50
51     if(argint(0, &n) < 0)
52         return -1;
53     addr = myproc()->sz;
54     myproc()->sz += n;
55
56     // if(growproc(n) < 0)
57     //     return -1;
58     return addr;
59 }
60
```

2) trap.c - Here we add condition to detect page fault in the switch section. As when a page fault occurs **tf->trapno** equals **T_PGFLT**. Here to handle this we added a function **PageFaulthandle()**. This function handles all the requests for page faults. **kalloc()** and **mappages()** were used to allocate required space or to inform user that no more memory is available.

```
case T_PGFLT:
    if(PageFaulthandle())<0){
        cprintf("Memory Not Allocated!!!");
    }
    break;
```

```

int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm);

int PageFaultHandle(){
    int addr1=rcr2();
    int addr2 = PGROUNDDOWN(addr1);
    char *mem=kalloc();
    if(mem!=0){
        memset(mem, 0, PGSIZE);
        if(mappages(myproc()->pgdir, (char*)addr2, PGSIZE, V2P(mem), PTE_W|PTE_U)<0)
            return -1;
        return 0;
    } else
        return -1;
}

```

3) **vm.c** - Return type of **mappages()** was changed from static int to int so that it can be used in **trap.c**

```

int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
{
    char *a, *last;
    pte_t *pte;

```

Answers to some Question

Q1 How does the kernel know which physical pages are used and unused?

Ans - kernel maintains the list of free pages in **kalloc.c** called **kmem**.

Q2 What data structures are used to answer this question?

Ans - A linked list is used as data structure for storing free pages.

Q3 Where do these reside?

Ans - Linked List is declared inside **kalloc.c** inside structure **kmem**.

Q4 Does xv6 memory mechanism limit the number of user processes?

Ans - Number of user processes are limited to 64 as defined by **NPROC** in **param.h**

Q5 If so, what is the lowest number of processes xv6 can 'have' at the same time?

Ans - Minimum number of processes in xv6 can be 1 as while starting only one process named **initproc** which initiates all other user processes.

Part B

Task 1

The **create_kernel_process()** function was created in **proc.c**. It took the name of the process and entrypoint function as arguments. It always remains in kernel mode. The parent process was set to **initproc** and **p->context->eip** was set to entrypoint argument and all other values as default. **allocproc** allocates the process a spot in the ptable. **setupkvm** maps the virtual address to physical address (from 0 to PHYSTOP).

```
void create_kernel_process(const char *name, void (*entrypoint)())
{
    struct proc *kp = allocproc();
    if(kp == 0){
        panic("Failed to allocate kernel process.");
    }
    kp->pgdir = setupkvm();
    if(kp->pgdir == 0)
    {
        kfree(kp->kstack);
        kp->kstack = 0;
        kp->state = UNUSED;
        panic("Failed to setup pgdir for kernel process.");
    }
    kp->sz = PGSIZE;
    kp->parent = initproc;
    memset(kp->tf, 0, sizeof(*kp->tf));
    kp->tf->cs = (SEG_KCODE << 3) | DPL_USER;
    kp->tf->ds = (SEG_KDATA << 3) | DPL_USER;
    kp->tf->es = kp->tf->ds;
    kp->tf->ss = kp->tf->ds;
    kp->tf->eflags = FL_IF;
    kp->tf->esp = PGSIZE;
    kp->tf->eip = 0;
    kp->tf->eax = 0;
    kp->cwd = namei("/");

    safestrcpy(kp->name, name, sizeof(name));

    acquire(&ptable.lock);

    kp->context->eip = (uint)entrypoint;
    kp->state = RUNNABLE;

    release(&ptable.lock);

    return;
}
```

Task 2

We create a container to contain the processes who have asked for additional memory but do not have any free pages. Therefore, we implement a circular queue and functions to give entry (cq_push) and exit (cq_pop) from the circular queue in proc.c .

```
166     struct circular_queue{
167         struct spinlock lock;
168         struct proc* queue[NPROC];
169         int head;
170         int tail;
171     };
172
173     // circular process queue for swapping out requests
174     struct circular_queue cq;
175
```

```
188
189     struct proc* cq_pop(){
190         acquire(&cq.lock);
191         if(cq.head == cq.tail){
192             release(&cq.lock);
193             return 0;
194         }
195         struct proc *p = cq.queue[cq.head];
196         cq.head = (cq.head + 1) % NPROC;
197         release(&cq.lock);
198         return p;
199     }
200
```

```
173
176     int cq_push(struct proc *p){
177         acquire(&cq.lock);
178         if ((cq.tail + 1) % NPROC == cq.head){
179             release(&cq.lock);
180             return 0;
181         }
182         cq.queue[cq.tail] = p;
183         cq.tail = (cq.tail + 1) % NPROC;
184         release(&cq.lock);
185
186         return 1;
187     }
188
```

We initialise the queue in userinit (user initialisation) function and lock for the queue in pinit function.

```
420
421 void
422 userinit(void)
423 {
424     acquire(&cq.lock);
425     cq.head = 0;
426     cq.tail = 0;
427     release(&cq.lock);
428
429     struct proc *p;
```

```
281 void
282 pinit(void)
283 {
284     initlock(&ptable.lock, "ptable");
285     initlock(&cq.lock, "cq");
286 }
```

We want to use the circular queue globally therefore we give its definition in defs.h

```
11 struct superblock;
12 struct circular_queue;
13
124 extern struct circular_queue cq;
125 int cq_push(struct proc *p);
126 int cp_pop();
```

Whenever a process needs to access some data it calls the **walkpgdir** function: If the data is not present in main memory, then **growproc** function is called which calls the **allocuvm** function which ultimately calls the **kalloc** function. If any free page is available then kalloc assigns it to the process else we need to swap out a page according to LRU policy to get a free page.

```
240 if(mem == 0){
241     // cprintf("allocuvm out of memory\n");
242     deallocuvm(pgdir, newsz, oldsz);
243
244     // SLEEP
245     myproc()->state = SLEEPING;
246     acquire(&sleeping_channel_lock);
247     myproc()->chan=sleeping_channel;
248     sleeping_channel_count++;
249     release(&sleeping_channel_lock);
250
251     cq_push(myproc());
252     if(!swap_out_process_exists){
253         swap_out_process_exists = 1;
254         create_kernel_process("swap_out_process", &swap_out_process_function);
255     }
256     return 0;
257 }
```

To swap out the page, we first need to move the process in the sleeping state on a special channel called **sleeping_channel**. We create this special channel in vm.c

```
14 struct spinlock sleeping_channel_lock;
15 int sleeping_channel_count = 0;
16 char *sleeping_channel;
17
18 // Set up CPU's kernel segment descriptors.
```

Then we declare it in defs.h to use it globally.

```
194 void clearpteu(pde_t *pgdir, char *uva);
195 extern struct spinlock sleeping_channel_lock;
196 extern int sleeping_channel_count;
197 extern char* sleeping_channel;
198
```

When we have a free page (either already had or after swap out) we need to assign it to the process, for this processes sleeping on sleeping_channel need to be woken up by **wakeup()** system call.

```
76 release(&kmem.lock);
77
78 // wake up processes sleeping on a sleeping channel
79 if(kmem.use_lock)
80     acquire(&sleeping_channel_lock);
81 if(sleeping_channel_count){
82     wakeup(sleeping_channel);
83     sleeping_channel_count = 0;
84 }
85 if(kmem.use_lock)
86     release(&sleeping_channel_lock);
87 }
```

We now implement the swap_out_process to really swap out the page.

To determine the victim page using LRU policy, we iterate through each entry in the process page table and look at the accessed bit which is obtained by bitwise & of the entry and PTE_A. The access bit indicates whether the page was accessed in the last iteration or not.

```
97 #define PTE_PS 0x080 // Page Size
98 #define PTE_A 0x020 // Accessed
99
```

```

231 void swap_out_process_function(){
232     acquire(&cq.lock);
233     while (cq.head != cq.tail){
234         struct proc *p = cq_pop();
235
236         pde_t *pd = p->pgdir;
237         for(int i = 0; i < NPENTRIES; i++){
238
239             // skip page table if accessed
240             if(pd[i] & PTE_A)
241                 continue;
242             pte_t *pt = (pte_t *) P2V(PTE_ADDR(pd[i]));
243             for(int j = 0; j < NPTENTRIES; j++){
244                 // skip if found
245                 if((pt[j] & PTE_A) || !(pt[j] & PTE_P))
246                     continue;
247                 pte_t *pte = (pte_t *) P2V(PTE_ADDR(pt[j]));
248
249                 // for file name
250                 int pid = p->pid;
251                 int virt = ((1 << 22) * i) + ((1 << 12) * j);
252
253                 // file name
254                 char c[50];
255                 itoa(pid, c);
256                 int x = strlen(c);
257                 c[x] = '_';
258                 itoa(virt, c + x + 1);
259                 safestrcpy(c + strlen(c), ".swp", 5);
260
261                 // file management
262                 int fd = proc_open(c, O_CREATE | O_RDWR);
263                 if (fd < 0){
264                     cprintf("Error creating or opening file: %s\n", c);
265                     panic("swap out process");
266                 }
267
268                 if(proc_write(fd, (char *) pte, PGSIZE) != PGSIZE){
269                     cprintf("Error writing to file: %s\n", c);
270                     panic("swap out process");
271                 }
272                 proc_close(fd);
273
274                 kfree((char *) pte);
275                 memset(&pt[j], 0, sizeof(pt[j]));
276
277                 // mark this page as swapped out
278                 pt[j] = pt[j] ^ 0x008;
279
280                 break;
281             }
282         }
283     }
284     release(&cq.lock);

```



```

254     release(&cq.lock);
255
256     struct proc *p;
257     if ((p = myproc()) == 0)
258         panic("swap out process");
259
260     swap_out_process_exists = 0;
261     p->parent = 0;
262     p->name[0] = '*';
263     p->killed = 0;
264     p->state = UNUSED;
265     sched();
266 }

```

In the scheduler, we unset the accessed bit.

```

654     for(int i = 0; i < NPENTRIES; i++){
655         // if PDE were accessed
656
657         if(((p->pgdir)[i] & PTE_P && ((p->pgdir)[i] & PTE_A){
658
659             pte_t *pt = (pte_t *) P2V(PTE_ADDR((p->pgdir)[i]));
660
661             for(int j = 0; j < NPTENTRIES; j++){
662                 if(pt[j] & PTE_A){
663                     pt[j] ^= PTE_A;
664                 }
665             }
666             ((p->pgdir)[i]) ^= PTE_A;
667         }
668     }
669 }

```

Now the swapped out page needs to be stored(written) in secondary storage for that we have copied open, write, close, functions from proc.c to sysfile.c and named as proc_open, proc_write, proc_close

```

8     #include "spinlock.h"
9     #include "fcntl.h"
10    #include "stat.h"
11    #include "sleeplock.h"
12    #include "fs.h"
13    #include "file.h"

```

, etc.


```

17  int
18  proc_close(int fd)
19  {
20      struct file *f;
21
22      if(fd < 0 || fd >= NOFILE || (f = myproc()->ofile[fd]) == 0)
23          return -1;
24      myproc()->ofile[fd] = 0;
25      fileclose(f);
26      return 0;
27  }
28
29  int
30  proc_write(int fd, char *p, int n)
31  {
32      struct file *f;
33
34      if(fd < 0 || fd >= NOFILE || (f = myproc()->ofile[fd]) == 0)
35          return -1;
36      return filewrite(f, p, n);
37  }

```

We clear the stack of the kernel process while exiting from it.

```

641      for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
642
643          // if the swapped out process has stopped running, free its stack
644          if(p->state == UNUSED && p->name[0] == '*'){
645              kfree(p->kstack);
646              p->kstack = 0;
647              p->name[0] = 0;
648              p->pid = 0;
649          }
650      }

```

Task-3

When a process had requested a data which is not in main memory, i.e, in case of a page fault, we need to swap the page containing data into the main memory.

First we need to create a container to satisfy the swap-in requests, just like we did in task-2. We create a circular queue and functions to enter and exit the container in proc.c.

```

struct circular_queue cq1;

int cq_push1(struct proc *p){
    acquire(&cq1.lock);
    if ((cq1.tail + 1) % NPROC == cq1.head){
        release(&cq1.lock);
        return 0;
    }
    cq1.queue[cq1.tail] = p;
    cq1.tail = (cq1.tail + 1) % NPROC;
    release(&cq1.lock);

    return 1;
}

struct proc* cq_pop1(){
    acquire(&cq1.lock);
    if(cq1.head == cq1.tail){
        release(&cq1.lock);
        return 0;
    }
    struct proc *p = cq1.queue[cq1.head];
    cq1.head = (cq1.head + 1) % NPROC;
    release(&cq1.lock);
    return p;
}

```

We initialise the queue in userinit (user initialisation) function and lock for the queue in pinit function.

```

void
userinit(void)
{
    acquire(&cq.lock);
    cq.head = 0;
    cq.tail = 0;
    release(&cq.lock);

    acquire(&cq1.lock);
    cq1.head = 0;
    cq1.tail = 0;
    release(&cq1.lock);
}

```

```

void
pinit(void)
{
    initlock(&ptable.lock, "ptable");
    initlock(&cq.lock, "cq");
    initlock(&sleeping_channel_lock, "sleeping_channel");
    initlock(&cq1.lock, "cq1");
}

```

We want to use the circular queue globally therefore we give its definition in defs.h

```

extern struct circular_queue cq1;
int cq_push1(struct proc *p);
struct proc* cq_pop1();

```

We create an integer variable to store the virtual address where the page fault has occurred in proc.h .

```

char name[16];           // Process name (debugging)
int va;                  //Virtual Address of the process

```

Whenever a page fault occurs the process traps the os, therefore to handle the page fault we add the following in trap.c

```

    break;
case T_PGFLT:
    pfhandling();
    break;

```

In pfhandling function, we set the process in sleeping state and obtain the virtual address where the page fault has occurred. Then we check whether the page was swapped out or not. If not then allow the default way of handling page fault else call **swap_in()** function.

```

19 void pfhandling(){
20     struct proc *p=myproc();
21     int va = rcr2();
22     acquire(&swaplock);
23     sleep(p,&swaplock);
24     pde_t *pde = &(p->pgdir)[PDX(va)];
25     pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
26
27     if((pgtab[PTX(va)])&0x008){
28         p->va = va;
29         cq_push1(p);
30         if(swap_in_process_exists==0)
31         {
32             swap_in_process_exists=1;
33             create_kernel_process("swap_in_process", &swap_in);
34         }
35     } else {
36         exit();
37     }
38 }

```

Then we declare swap_in() in defs.h to use it globally.

```

void        swap_out_process_function();
void        swap_in();

extern int   swap_out_process_exists;
extern int   swap_in_process_exists;

```

We have already implemented other file management functions in task-2. We here implement read_process which is basically copy function.

```

int read_process(int fd, int n, char *p)
{
    struct file *fl;
    if(fd < 0 || fd >= NOFILE) return -1;
    fl =myproc()->ofile[fd];
    if(fl==0) return -1;
    return fileread(fl, p, n);
}

```

We allocate a free page to the process in main memory and read the page from secondary storage to the allocated free page in the main memory.

```
void swap_in(){
    acquire(&cq1.lock);
    while(cq1.head!=cq1.tail){
        struct proc *p=cq_pop1();
        int process_id=p->pid;
        int va=PTE_ADDR(p->va);
        char pagename[50];
        int_to_string(process_id,pagename);
        int length =strlen(pagename);
        pagename[length]='_';
        int_to_string(va,pagename+length+1);
        safestrcpy(pagename+strlen(pagename),".swp",5);

        int fd=proc_open(pagename,0_RDONLY);
        if(fd<0){
            release(&cq1.lock);
            cprintf("Page could not be found in memory: %s\n", pagename);
            panic("swap in failed");
        }
        char *mem=kalloc();
        read_process(fd,PGSIZE,mem);
        int x = PTE_W|PTE_U;
        int mp = mappages(p->pgdir, (void *)va, PGSIZE, V2P(mem),x );
        if(mp<0){
            release(&cq1.lock);
            panic("page mapping");
        }
        wakeup(p);
    }
    release(&cq1.lock);
    struct proc *p = myproc();
    if(p==0)
        panic("swap in failed");
    swap_in_process_exists=0;
    p->parent = 0;
    p->name[0] = '*';
    p->killed = 0;
    p->state = UNUSED;
    sched();
}
```


Task-4:Sanity Test

We will create a user-space program to test our swapping mechanism.

20 child processes are created using fork().

10 4Kb blocks of memory is allocated for each process.

For a given process_id <pid>, block number <j> and offset <k>, the memory location stored in address field is $pid * 100000 + j * 10000 + k$.

```
int main(int argc, char *argv[]){
    int *add_list[10];
    int id_list[100];
    int counter = 0;
    for (int i = 0; i < 10; i++){
        {
            if (fork() == 0){
                counter = counter + 1;
                for (int j = 0; j < 10; j++){
                    int *addr = (int *)malloc(4096);
                    int p_id;
                    if ((add_list[j] = addr) == NULL){
                        p_id = getpid();
                        printf(1, "the process ID is: %d\n", p_id);
                        break;
                    }
                    p_id = getpid();
                    id_list[counter] = p_id;
                    for (int k = 0; k < 1024; k++){
                        *(addr + k) = p_id * 100000 + j * 10000 + k;
                    }
                    if (j == 0)
                        printf(1, "block 1 index : %d, Beginning Address : %p , process ID : %d\n", counter, add_list[j], p_id);
                    if (j == 4)
                        printf(1, "block 5 index : %d, Beginning Address : %p , process ID : %d\n", counter, add_list[j], p_id);
                    if (j == 9)
                        printf(1, "block 10 index : %d, Beginning Address : %p , process ID : %d\n", counter, add_list[j], p_id);
                }
            }
            else break;
        }
    }
    while (wait() != -1); // Execute all the child process
    if (counter == 0)
        exit();
    for (int i = 0; i < 10; i++){
        if (i == 0)
            printf(1, "Beginning Address : %p ,Process ID: %d, 100th value of the 1st block: %d \n", add_list[i], id_list[counter], *(add_list[i] + 100));
        if (i == 4)
            printf(1, "Beginning Address : %p ,Process ID: %d, 100th value of the 5th block: %d \n", add_list[i], id_list[counter], *(add_list[i] + 100));
        if (i == 9)
            printf(1, "Beginning Address : %p ,Process ID: %d, 100th value of the 10th block: %d \n", add_list[i], id_list[counter], *(add_list[i] + 100));
    }
    counter--;
    exit();
}
```

Output:

- 1) Details of block 1, 5, 10 are displayed on the console.
- 2) After all child processes stop executing the contents of memory locations are checked.

```
$ sanity
block 1 index : 1, Beginning Address : A000 , process ID : 4
block 5 index : 1, Beginning Address : 5FE0 , process ID : 4
block 10 index : 1, Beginning Address : FFF0 , process ID : 4
block 1 index : 2, Beginning Address : EFE8 , process ID : 5
block 5 index : 2, Beginning Address : 1A000 , process ID : 5
block 10 index : 2, Beginning Address : 14FD8 , process ID : 5
block 1 index : 3, Beginning Address : 13FD0 , process ID : 6
block 5 index : 3, Beginning Address : 1EFE8 , process ID : 6
block 10 index : 3, Beginning Address : 28FF8 , process ID : 6
block 1 index : 4, Beginning Address : 27FF0 , process ID : 7
block 5 index : 4, Beginning Address : 23FD0 , process ID : 7
block 10 index : 4, Beginning Address : 2DFE0 , process ID : 7
block 1 index : 5, Beginning Address : 2CFD8 , process ID : 8
block 5 index : 5, Beginning Address : 37FF0 , process ID : 8
block 10 index : 5, Beginning Address : 42000 , process ID : 8
Beginning Address : 2CFD8 ,Process ID: 8, 100th value of the 1st block: 800100
Beginning Address : 37FF0 ,Process ID: 8, 100th value of the 5th block: 840100
Beginning Address : 42000 ,Process ID: 8, 100th value of the 10th block: 890100
Beginning Address : 27FF0 ,Process ID: 7, 100th value of the 1st block: 700100
Beginning Address : 23FD0 ,Process ID: 7, 100th value of the 5th block: 740100
Beginning Address : 2DFE0 ,Process ID: 7, 100th value of the 10th block: 790100
Beginning Address : 13FD0 ,Process ID: 6, 100th value of the 1st block: 600100
Beginning Address : 1EFE8 ,Process ID: 6, 100th value of the 5th block: 640100
Beginning Address : 28FF8 ,Process ID: 6, 100th value of the 10th block: 690100
Beginning Address : EFE8 ,Process ID: 5, 100th value of the 1st block: 500100
Beginning Address : 1A000 ,Process ID: 5, 100th value of the 5th block: 540100
Beginning Address : 14FD8 ,Process ID: 5, 100th value of the 10th block: 590100
Beginning Address : A000 ,Process ID: 4, 100th value of the 1st block: 400100
Beginning Address : 5FE0 ,Process ID: 4, 100th value of the 5th block: 440100
Beginning Address : FFF0 ,Process ID: 4, 100th value of the 10th block: 490100
$
```

0xE000000

Results:-

- 1) When phystop=0xE000000, then we have all 20 processes in the output.
- 2) When phystop=0x0400000, then we have only 9 out of 20 processes in the output.

Explanation:-When we reduce the phystop then, we don't have enough capacity to hold all the processes in the main memory.

NOTE:-

- 1) Part-A and Part-B are separately implemented on two different xv-6 directories.
- 2) Use <patch -ruN -strip -d xv6-public < patch_A.txt> to apply patch for A part.
- 3) Use <patch -ruN -strip -d xv6-public < patch_B.txt> to apply patch for B part.